

VSAR Inference Engine: Unified Multi-Mode Control Semantics

This document specifies a **complete inference engine** for VSAR that exercises all supported reasoning modes using the encodings defined in “**VSAR: A Unified Encoding & Reasoning Framework (Clean Edition)**”.

VSAR’s core idea:

One representation substrate (FHRR encodings) + multiple inference controllers (semantics).

This engine is built as a **modular, pluggable control layer** over a shared memory of encoded objects.

0) Summary of reasoning modes and corresponding controllers

Mode	Controller	Primary operations	Output type
Deductive / Horn	Forward/backward chaining	unify, apply rule, index	derived facts
DL (ALC core)	Tableau / completion	expand, clash check	satisfiable? / entailment
Abduction	Goal-driven inverse rules	unify head, hypothesize body	explanations
Induction	Program search / rule learning	generalize, score	learned rules
Analogical	Structure mapping	retrieve, align, map	mappings + transfers
Case-based	kNN retrieval + adaptation	similarity, mapping	adapted solution
Probabilistic	Weighted propagation	t-norms, noisy-OR	belief strengths
Paraconsistent	4-valued / non-explosive	separate supports	$\langle \text{supp}(A), \text{supp}(\neg A) \rangle$
Non-monotonic/default	Defeasible + defeat	priorities, exceptions	warranted set
Argumentation	AF/BAF solver	support/attack, propagation	acceptability scores

Mode	Controller	Primary operations	Output type
Epistemic	Multi-agent belief models	agent-indexed KBs	K/B entailments

All controllers share:

- the same encodings and typed codebooks
- the same unification kernel
- the same retrieval/cleanup infrastructure

1) Data model and memory layout

1.1 Core stores

VSAR maintains these stores (conceptually; implementation may vary):

- **Symbol memories** (typed codebooks): E, C, R, F, P, TAGS, OPS, ROLES...
- **Fact store**: set of stored ground literals
- **Rule store**: stored rules (Horn, defeasible, DL axioms)
- **Graph store**: argument graphs (support/attack edges)
- **Case store**: cases for CBR
- **Map store**: analogy mappings

1.2 Item schema (unified)

Every stored item is:

```
Item = {
    vec: Complex[d],
    kind: enum {FACT, RULE, AXIOM, EDGE, CASE, MAP},
    weight: float or interval or mass,
    priority: float (optional),
    provenance: {source, timestamp, agent, trace},
    tags: set,
}
```

1.3 Paraconsistent belief record

For each literal L , maintain independent supports:

```
Belief[L] = (supp_pos, supp_neg)
```

- supp_pos applies to L

- supp_neg applies to $\neg L$

This is the basis for 4-valued semantics.

2) Core primitives used by all controllers

2.1 Similarity, retrieval, and indexing

VSAR uses two *distinct* similarity mechanisms:

- 1) **Cleanup similarity (typed)**: nearest basis symbol in a codebook
- 2) **Item retrieval similarity (exploratory)**: nearest stored items to a query vector

(A) Cleanup

Given v and a typed codebook S :

$$\text{cleanup}_S(v) = \arg \max_{s \in S} \cos(v, s)$$

Cleanup is a **commitment** operation.

(B) Retrieval

Given query q and store \mathcal{M} :

$$\text{retrieve}_k(q) = \text{TopK}_{x \in \mathcal{M}} \cos(q, x.\text{vec})$$

Retrieval is **exploratory** and returns a set of candidates.

2.2 Unification kernel (slot-level)

Unification is performed by **structure-aware decoding** rather than whole-vector similarity.

2.2.1 Decode atom/term slots

For an atom encoding:

$$A = (P \otimes \text{TAG}_{ATOM}) \otimes \bigoplus_i (\text{ARG}_i \otimes \text{enc}(t_i))$$

Slot payload:

$$\text{payload} = A \oslash (P \otimes \text{TAG}_{ATOM})$$

Candidate slot value:

$$v_i = payload \oslash ARG_i$$

Then apply **typed cleanup**:

- constants: `cleanup_E(v_i)`
- terms: `cleanup_TERM(v_i)` (or recursively decode as term structure)

2.2.2 Variable handling

Variables are not stored as basis symbols in the entity codebook. Instead, represent a variable token as:

- `VAR ⊗ TAG_META ⊗ (ID ⊗ var_id)`

Unification returns a substitution map:

$$\theta = \{ x \mapsto a, y \mapsto f(b), \dots \}$$

Operationally, a variable position is decoded to a vector; the engine searches for bindings consistent with constraints.

2.3 Crisp query answering

For patterns like `p(a, ?)`:

1) retrieve candidate facts with matching predicate (index by predicate)
 2) for each fact, decode the unknown slot via unbind
 3) cleanup to a typed symbol
 4) return all cleaned answers above threshold

This avoids “fuzzy whole-fact similarity” and produces crisp results.

3) Deductive reasoning (Horn / Datalog-style)

3.1 Rule representation

Rules are stored as:

$$R = TAG_{RULE} \otimes (HEAD \otimes enc(H) \oplus BODY \otimes enc(B_1 \wedge \dots \wedge B_n))$$

3.2 Forward chaining algorithm (semi-naive)

Inputs: facts F_0 , rules \mathcal{R}

Loop:

1. Maintain delta facts Δ
2. For each rule, match body literals against $F \cup \Delta$ using unification

3. Produce new head instances and add to Δ'
4. Stop when Δ' empty

VSAR-specific optimization:

- index facts by predicate vector P_p
- use slot decoding + cleanup for constant checks
- batch unbinding and cleanup in JAX

3.3 Backward chaining (goal-driven)

Procedure prove(goal):

- retrieve rules whose head predicate matches goal predicate
- unify goal with rule head
- recursively prove each instantiated body literal

Use memoization (tabling) to avoid loops.

4) Description Logic reasoning (ALC core)

VSAR supports DL by storing:

- TBox axioms: $C \sqsubseteq D, C \equiv D$
- ABox assertions: $C(a), R(a, b)$

4.1 Completion / tableau state

Maintain a completion graph:

- node = individual
- $\text{label}(\text{node})$ = set of concept expressions
- edges = role assertions

4.2 Expansion rules (ALC)

- **n-rule:** if $C \sqcap D \in L(x)$ then add C, D to $L(x)$
- **u-rule:** if $C \sqcup D \in L(x)$ then branch: add C or add D
- **\exists -rule:** if $\exists R.C \in L(x)$ then create y with edge $R(x,y)$ and add C to $L(y)$
- **\forall -rule:** if $\forall R.C \in L(x)$ and $R(x,y)$ then add C to $L(y)$
- **clash:** if $C \in L(x)$ and $\neg C \in L(x)$ then clash

4.3 VSAR-specific machinery

- decode constructors by unbinding $\text{LEFT}/\text{RIGHT}/\text{ROLE}/\text{FILLER}$
- use typed cleanup to recognize operator tags
- retrieval is used to pull relevant axioms by matching LHS concept heads

4.4 Paraconsistent DL (optional)

If paraconsistent mode is enabled, a clash does **not** explode; it updates belief state and allows argumentation/defeat to resolve.

5) Paraconsistent inference (4-valued)

5.1 Belief update

When a fact L arrives with weight w :

- if L is positive: $\text{supp}(L) \leftarrow \text{combine}(\text{supp}(L), w)$
- if L is negative: $\text{supp}(\neg L) \leftarrow \text{combine}(\text{supp}(\neg L), w)$

Combine can be max, noisy-OR, or log-odds merge.

5.2 Query answering under paraconsistency

Return both supports:

```
query(L) -> (supp(L), supp(\neg L))
```

Downstream controllers interpret this as true/false/both/neither.

6) Non-monotonic and default reasoning

Defaults are defeasible rules with priorities and exceptions.

6.1 Defeasible rule structure

```
Rule = (vec, weight, priority, exceptions)
```

6.2 Defeat relation

A derived conclusion is **defeated** if there exists an attacking rule/argument with:

- higher priority OR
- equal priority but higher weight OR
- more specific conditions

6.3 Warrant computation

Compute a fixpoint of warranted conclusions using an argumentation solver (Section 8).

7) Probabilistic / weighted inference

All derived items carry weights.

7.1 Body combination (AND)

$$w_{body} = T(w_1, \dots, w_n)$$

where T is a pluggable t-norm (product, min, etc.).

7.2 Proof combination (OR)

$$w_{head} = S(w^{(1)}, \dots, w^{(m)})$$

where S is a pluggable t-conorm (max, noisy-OR, etc.).

7.3 Integration with paraconsistency

Maintain weights independently for L and $\neg L$.

8) Argumentation engine (AF/BAF)

8.1 Graph construction

- nodes: claims (literals, axioms, conclusions)
- edges: support and attack relations

Edge encodings:

$$enc(support(A, B)) = (EDGE \otimes SUPPORT) \otimes (SRC \otimes enc(A) \oplus TGT \otimes enc(B))$$

8.2 Semantics

Choose a semantics family:

- Dung AF (preferred extensions)
- bipolar AF (support + attack)
- weighted propagation (gradual semantics)

8.3 Gradual acceptability (example)

Let each node have base strength from evidence, then iteratively update:

$$str_{t+1}(x) = \sigma \left(base(x) + \sum_{s \in Sup(x)} \alpha str_t(s) - \sum_{a \in Att(x)} \beta str_t(a) \right)$$

where σ is a squashing function.

Warranted conclusions are those with acceptability above threshold.

9) Epistemic reasoning (multi-agent)

9.1 Agent-indexed KBs

Maintain separate stores per agent:

- `KB[a].facts`
- `KB[a].rules`
- `KB[a].beliefs`

9.2 Modal operators

Knowledge/belief objects are encoded and can be nested:

- `enc(K_a φ)`
- `enc(B_a φ)`

9.3 Reasoning policies

- **K entailment** (factive, optional): if `K_a φ` then ϕ in world model
- **B entailment**: if `B_a φ` then ϕ in $KB[a]$ with some weight

Updates:

- observation update to world
- communication update across agents (with trust weights)

Paraconsistency is supported per-agent.

10) Abduction engine

Given a goal literal G:

1) retrieve rules with head predicate matching G 2) unify G with rule head to obtain substitution θ 3) instantiate body under θ 4) identify which body literals are missing in current KB 5) propose missing as hypotheses

10.1 Explanation scoring

Score an explanation set H by:

$$Score(H) = \lambda_1 \sum_{h \in H} \log P(h) + \lambda_2 \sum_{r \in used} \log w(r) - \lambda_3 |H| - \lambda_4 Clash(H, KB)$$

Optionally include argumentation acceptability.

Return top-k explanations.

11) Analogical reasoning engine

11.1 Retrieval of candidate structures

- retrieve similar relational structures (facts/rules/cases) by exploratory similarity
- decode structures into slot-level components

11.2 Structure mapping

Compute a mapping M between symbols in domain A and domain B that maximizes preserved relations.

One objective:

$$M^* = \arg \max_M \sum_{(p(a,b) \in A)} \mathbf{1}[p(M(a), M(b)) \in B] \cdot w(p)$$

Use resonators / constraint satisfaction to find consistent M.

11.3 Transfer

Apply mapping to produce candidate hypotheses or solutions.

12) Case-based reasoning engine

Given a new problem P:

- 1) encode as CASE(PROB \otimes enc(P) \oplus CONTEXT \otimes enc(ctx))
 - 2) retrieve top-k similar cases by similarity of PROB slot
 - 3) adapt solutions using analogical mapping engine
 - 4) store new case with outcome and provenance
-

13) Induction (rule learning)

Induction is implemented as a search over rule templates.

13.1 Candidate generation

- anti-unification over sets of facts
- generalize constants into variables
- propose Horn clause templates

13.2 Scoring

Score candidate rules by:

- coverage of positive examples
- rejection of negatives
- simplicity
- weight/probability calibration

Store learned rules as defeasible rules with provenance.

14) Orchestration: a single engine that exercises all modes

14.1 Core loop

On each update (new facts/axioms/cases):

1) **Update belief state** (paraconsistent supports) 2) **Run deductive closure** (Horn) under current warranted set 3) **Run DL completion** for relevant individuals/concepts 4) **Recompute argument graph** and acceptability 5) **Apply defaults** (defeasible) subject to defeat

On each query:

- if it is a crisp logical query: use unbind→cleanup and deductive/proof controllers
- if it is an entailment query: consult DL + Horn + argumentation
- if it is an explanation query: invoke abduction
- if it is an analogy query: invoke mapping
- if it is a similar-case query: invoke CBR
- if it is agent-indexed: invoke epistemic controller

14.2 Guarantees and intended behavior

- **Soundness** is logic-dependent (DL/Horn when run classically)
 - **Robustness** is provided by VSA similarity and noise tolerance
 - **Non-explosiveness** is guaranteed by paraconsistent belief tracking
 - **Extensibility** comes from adding new controllers without changing encodings
-

15) Engine-level test plan (high-level)

- deterministic bind/unbind invariants
 - cleanup correctness under controlled bundling noise
 - unification correctness on nested terms
 - Horn closure matches symbolic baseline on toy KB
 - ALC tableau results match a reference reasoner on toy ontologies
 - paraconsistent non-explosion: derive does not trivialize under A and $\neg A$
 - default defeat behaves by priority
 - abductive explanations recover known hidden causes
 - argumentation acceptability aligns with chosen semantics
 - epistemic nesting decoding and agent separation
 - analogy transfers preserve relational structure
 - CBR retrieval/adaptation correctness
-

16) Implementation notes (VSAX-optimized)

- batch all unbind operations across candidate facts/rules
 - use typed codebooks with matrix multiplication for cleanup
 - maintain predicate indexes for O(1) candidate narrowing
 - keep separate stores per agent for epistemic reasoning
 - represent all proofs/derivations as provenance graphs for interpretability
-

17) Final view

VSAR's inference engine is a **portfolio of semantics** running over a single representational layer.

You can start with a minimal subset (Horn + cleanup) and progressively add controllers without changing encodings.